

e-feature-maps-of-a-cnn-model-lab8

May 8, 2023

[1]: `!pip install torchsummary`

```
Collecting torchsummary
  Downloading torchsummary-1.5.1-py3-none-any.whl (2.8 kB)
Installing collected packages: torchsummary
Successfully installed torchsummary-1.5.1
WARNING: Running pip as the 'root' user can result in broken permissions
and conflicting behaviour with the system package manager. It is recommended to
use a virtual environment instead: https://pip.pypa.io/warnings/venv
```

[2]: `# import libraries`

```
import os
import sys
import time

import numpy as np
import pandas as pd

import matplotlib.pyplot as plt

from sklearn.metrics import accuracy_score

import torch
import torch.nn as nn
from torch.utils.data import DataLoader
from torchvision.datasets import ImageFolder
import torchvision.transforms as T
from torchsummary import summary
```

[3]: `# select GPU if available`

```
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
# device = "cpu" # explicitly set to execute on CPU
```

[4]: `!nvidia-smi`

```
Mon May  8 17:26:36 2023
```

NVIDIA-SMI 470.161.03 Driver Version: 470.161.03 CUDA Version: 11.4							
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr.	ECC
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute M.	MIG M.
0	Tesla P100-PCIE...	Off	00000000:00:04.0	Off			0
N/A	37C	P0	26W / 250W	2MiB / 16280MiB	0%	Default	N/A

Processes:							
GPU	GI	CI	PID	Type	Process name	GPU Memory	Usage
ID	ID						
No running processes found							

```
[5]: # transformation of the training/testing images
transformer = T.Compose([
    T.Resize((150,150)),
    T.RandomHorizontalFlip(),
    T.ToTensor(), #0-255 to 0-1, numpy to tensors
    T.Normalize([0.5,0.5,0.5], # 0-1 to [-1,1] , formula (x-mean)/std
               [0.5,0.5,0.5])
])

train_path = "/kaggle/input/intel-image-classification/seg_train/seg_train/"
test_path = "/kaggle/input/intel-image-classification/seg_test/seg_test/"

# load the training/testing images
train_set = ImageFolder(train_path, transform = transformer)
test_set = ImageFolder(test_path, transform = transformer)

# create DataLoader
batch_size = 256
train_loader = DataLoader(train_set, batch_size = batch_size, shuffle = True)
test_loader = DataLoader(test_set, batch_size = batch_size, shuffle = True)

[6]: # total test images, total train images
len(test_set), len(train_set)
```

```
[6]: (3000, 14034)
```

```
[7]: # categories and their labels  
train_set.class_to_idx
```

```
[7]: {'buildings': 0,  
      'forest': 1,  
      'glacier': 2,  
      'mountain': 3,  
      'sea': 4,  
      'street': 5}
```

```
[8]: # create model architecture  
class CNN_Net(nn.Module):  
    def __init__(self, is_print_toggle = False):  
        super(CNN_Net, self).__init__()  
  
        ##-----convolution  
        # layers-----  
  
        ## 1st convolution - max_pooling block ==> # np.floor((image_height +  
        ↵2*padding - kernel_size)/stride) + 1  
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=12, kernel_size=3, u  
        ↵stride=1, padding=1)  
        # shape = (256, 12, 150, 150)  
        self.bn1 = nn.BatchNorm2d(num_features=12)  
        # shape = (256, 12, 150, 150)  
        self.relu1 = nn.ReLU()  
        # shape = (256, 12, 150, 150)  
        self.max_pool1 = nn.MaxPool2d(kernel_size = 2)  
        # So, after maxpool, 148/2 = 74 # shape = (256, 12, 75, 75)  
  
        ## 2nd convolution - max_pooling block  
        self.conv2 = nn.  
        ↵Conv2d(in_channels=12,out_channels=20,kernel_size=3,stride=1,padding=1)  
        # shape = (256, 20, 75, 75)  
        self.relu2 = nn.ReLU()  
        # shape = (256, 20, 75, 75)  
  
        ## convolution layer 3  
        self.conv3 = nn.Conv2d(in_channels=20,u  
        ↵out_channels=32,kernel_size=3,stride=1,padding=1)  
        # shape = (256, 32, 75, 75)  
        self.bn3 = nn.BatchNorm2d(num_features=32)  
        # shape = (256, 32, 75, 75)
```

```

    self.relu3      = nn.ReLU()
    # shape = (256, 32, 75, 75)

    ##-----FCL-----
    estimated_features = 32 * 75 * 75 # output_channel * height * width = 180000 (from convolution layer 3)

    # output layer
    self.output = nn.Linear(in_features = estimated_features, out_features= 6)

    self.print_toggle = is_print_toggle

def forward(self, X):

    ##-----convolution-----
    layers
    print(f"Input: {X.shape}") if self.print_toggle else None

    # 1st convolution-maxpooling block
    convolution_layer1 = self.relu1( self.bn1( self.conv1(X) ) )
    X = self.max_pool1( convolution_layer1 )
    print(f"conv1/max_pool: {X.shape}") if self.print_toggle else None

    # 2nd convolution-maxpooling block
    convolution_layer2 = self.relu2( self.conv2( X ) )
    print(f"Layer conv2/max_pool: {X.shape}") if self.print_toggle else None

    # 3rd convolution-maxpooling block
    X = self.relu3( self.bn3( self.conv3( convolution_layer2 ) ) )
    print(f"Layer conv3/max_pool: {X.shape}") if self.print_toggle else None

    ##-----FCL-----
    # reshape for linear layers
    nUnits = X.shape.numel()/X.shape[0]
    X = X.view(-1, int(nUnits)) # this shape should be (256, 32 * 75 * 75) = (256, 180000)
    if self.print_toggle: print(f"Vectorize: {X.shape}")

```

```

# Output layer
X = self.output(X)
if self.print_toggle: print(f"Layer output: {X.shape}")

return X, convolution_layer1, convolution_layer2

```

```
[9]: # create neural network
def create_network(is_print_toggle, learning_rate = 0.001):
    # create an model instance
    network = CNN_Net(is_print_toggle)

    # select loss function
    loss_fun = nn.CrossEntropyLoss()

    # select optimizer
    optimizer = torch.optim.Adam(network.parameters(), lr = learning_rate,
                                weight_decay=0.0001)

    return network, loss_fun, optimizer

```

```
[10]: # test the model with one batch (sanity check)
network, loss_fun, optimizer = create_network(is_print_toggle = True)

X, y = iter(train_loader).next()
yHat = network(X)[0]  # here, we dont need the convolution layers' outputs

# check sizes of model outputs and target variable
print(' ')
print(yHat.shape)
print(y.shape)

# now let's compute the loss
loss = loss_fun(yHat, y)
print(' ')
print('Loss: ')
print(loss)
```

```

Input: torch.Size([256, 3, 150, 150])
conv1/max_pool: torch.Size([256, 12, 75, 75])
Layer conv2/max_pool: torch.Size([256, 12, 75, 75])
Layer conv3/max_pool: torch.Size([256, 32, 75, 75])
Vectorize: torch.Size([256, 180000])
Layer output: torch.Size([256, 6])

torch.Size([256, 6])
torch.Size([256])

```

```
Loss:  
tensor(1.8292, grad_fn=<NllLossBackward0>)
```

```
[11]: # summary of the model  
network, loss_fun, optimizer = create_network(is_print_toggle = False)  
  
if device == "cpu":  
    summary(network.to(device), (3, 150, 150), device=device);  
else:  
    summary(network.to(device), (3, 150, 150));
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 12, 150, 150]	336
BatchNorm2d-2	[-1, 12, 150, 150]	24
ReLU-3	[-1, 12, 150, 150]	0
MaxPool2d-4	[-1, 12, 75, 75]	0
Conv2d-5	[-1, 20, 75, 75]	2,180
ReLU-6	[-1, 20, 75, 75]	0
Conv2d-7	[-1, 32, 75, 75]	5,792
BatchNorm2d-8	[-1, 32, 75, 75]	64
ReLU-9	[-1, 32, 75, 75]	0
Linear-10	[-1, 6]	1,080,006

Total params: 1,088,402

Trainable params: 1,088,402

Non-trainable params: 0

Input size (MB): 0.26

Forward/backward pass size (MB): 12.53

Params size (MB): 4.15

Estimated Total Size (MB): 16.94

```
[12]: # create a function to train the model  
def train_the_model(network, loss_fun, optimizer, epochs = 100):  
    # number of epochs  
    epochs = epochs  
  
    # send the model to GPU  
    network.to(device)  
  
    # initialize losses  
    train_loss      = torch.zeros(epochs)  
    test_loss       = torch.zeros(epochs)
```

```

train_accuracy = []
test_accuracy  = []

# loop over epochs
for epoch in range(epochs):
    # To calculate the execution time,
    start_time = time.time()

    # switch on the training mode
    network.train()

    batch_accuracy = []
    batch_loss      = []
    # loop over the training data batches
    for X, y in train_loader:
        # push X and y to GPU
        X = X.to(device)
        y = y.to(device)

        # forward propagation
        yHat = network(X)[0] # here, we dont need the convolution layers'
↳outputs

        # compute loss
        loss = loss_fun(yHat, y)

        # backward propagation
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # store the loss from this batch
        batch_loss.append(loss.item())

        # compute accuracy
        accuracy_percentage = accuracy_score(y.detach().cpu(), torch.
↳argmax(yHat, axis=1).detach().cpu()) * 100
        batch_accuracy.append(accuracy_percentage)
# end of batch loop

# now that we've trained through the batches, get their average train
↳accuracy
train_acc = np.mean(batch_accuracy)
train_accuracy.append( train_acc )

# and get average losses across the batches
train_loss[epoch] = np.mean( batch_loss )

```

```

# switch to evaluation mode to calculate test accuracy
network.eval()

# # extract X,y from test dataloader
# X, y = next(iter(test_loader))

temp_test_loss = [] # to store loss temporarily
temp_test_acc = [] # to store accuracy temporarily
# loop through the test batches
for X, y in test_loader:
    # push the data to GPU
    X = X.to(device)
    y = y.to(device)

    with torch.no_grad(): # deactivate autograd
        # calculate prediction
        yHat = network(X)[0] # here, we dont need the convolution
        ↪ layers' outputs
        # calculate loss
        loss = loss_fun(yHat, y)
        # and store it
        temp_test_loss.append(loss.item())

        # calculate accuracy and store it
        accuracy_percentage = accuracy_score(y.detach().cpu(), torch.
        ↪ argmax(yHat, axis=1).detach().cpu()) * 100
        temp_test_acc.append(accuracy_percentage)

    # calculate overall test loss and store it
    test_loss[epoch] = np.mean(temp_test_loss)

    # calculate the test accuracy
    test_acc = np.mean(temp_test_acc)
    test_accuracy.append(test_acc)

    # calculate end time
    end_time = time.time()

    # display a status message
    msg = f"Completed: {epoch + 1}/{epochs} - took: { (end_time - start_time) :.6f} sec - Train accuracy: {train_acc:.5f} - Test accuracy:{test_acc:.5f}"
    sys.stdout.write("\n" + msg)

# end epochs

```

```
# function output
return train_loss, test_loss, train_accuracy, test_accuracy, network
```

[13]: %%time

```
# create a neural network
network, loss_fun, optimizer = create_network(is_print_toggle = False)
# train the model
train_loss, test_loss, train_accuracy, test_accuracy, network = train_the_model(network, loss_fun, optimizer, epochs = 10)
```

```
Completed: 1/10 - took: 88.125844 sec - Train accuracy: 52.29742 - Test
accuracy: 53.71518
Completed: 2/10 - took: 43.782111 sec - Train accuracy: 69.92796 - Test
accuracy: 66.42889
Completed: 3/10 - took: 44.413203 sec - Train accuracy: 78.23410 - Test
accuracy: 71.79857
Completed: 4/10 - took: 44.385618 sec - Train accuracy: 85.71659 - Test
accuracy: 72.18354
Completed: 5/10 - took: 44.781815 sec - Train accuracy: 86.47497 - Test
accuracy: 63.52327
Completed: 6/10 - took: 44.121289 sec - Train accuracy: 90.64516 - Test
accuracy: 73.69933
Completed: 7/10 - took: 44.227841 sec - Train accuracy: 93.64854 - Test
accuracy: 70.85598
Completed: 8/10 - took: 45.662771 sec - Train accuracy: 92.90192 - Test
accuracy: 68.60988
Completed: 9/10 - took: 44.171983 sec - Train accuracy: 95.42012 - Test
accuracy: 73.04688
Completed: 10/10 - took: 44.237915 sec - Train accuracy: 93.71199 - Test
accuracy: 74.64193CPU times: user 4min 56s, sys: 56 s, total: 5min 52s
Wall time: 8min 7s
```

[14]: # display the accuracy of our model

```
fig, ax = plt.subplots(2, 1, figsize = (10, 9))

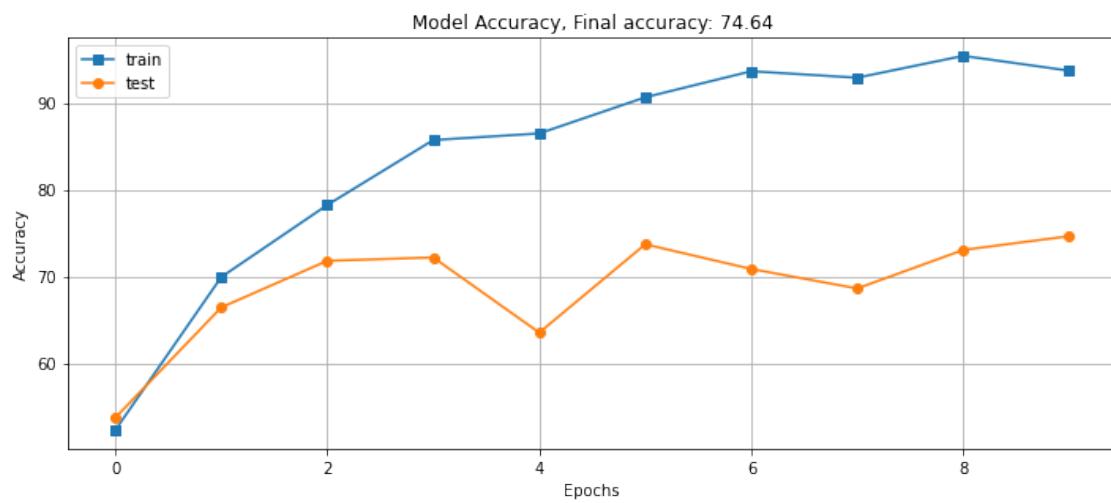
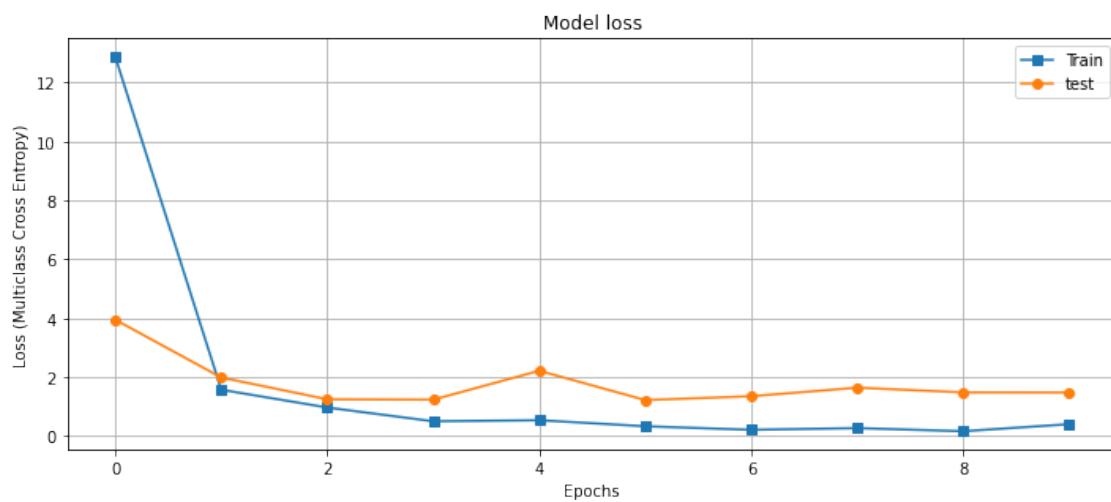
ax[0].plot(train_loss, "s-", label = "Train")
ax[0].plot(test_loss, "o-", label = "test")
ax[0].set_xlabel("Epochs")
ax[0].set_ylabel("Loss (Multiclass Cross Entropy)")
ax[0].set_title("Model loss")
ax[0].legend()
ax[0].grid()
```

```

ax[1].plot(train_accuracy, "s-", label = "train")
ax[1].plot(test_accuracy, "o-", label = "test")
ax[1].set_xlabel("Epochs")
ax[1].set_ylabel("Accuracy")
ax[1].set_title(f"Model Accuracy, Final accuracy: {test_accuracy[-1]:.2f}")
ax[1].legend()
ax[1].grid()

plt.tight_layout()
plt.show()

```



[15]: *### now let's look though the convolution channels ###*

```

# extract a batch from the test set
X, y = next(iter(test_loader))

```

```

# move the data to GPU
X = X.to(device)
y = y.to(device)

# prediction and out convolution layers
yHat, feature_maps1, feature_maps2= network(X)
# check the shape of the two feature maps
print(feature_maps1.shape, feature_maps2.shape)

```

torch.Size([256, 12, 150, 150]) torch.Size([256, 20, 75, 75])

[16]: # Feature maps from the convolution layer 1

```

number_of_images = 10

fig,axs = plt.subplots(feature_maps1.shape[1] + 1, number_of_images,□
    ↪ figsize=(70, 70))

for pici in range(number_of_images):

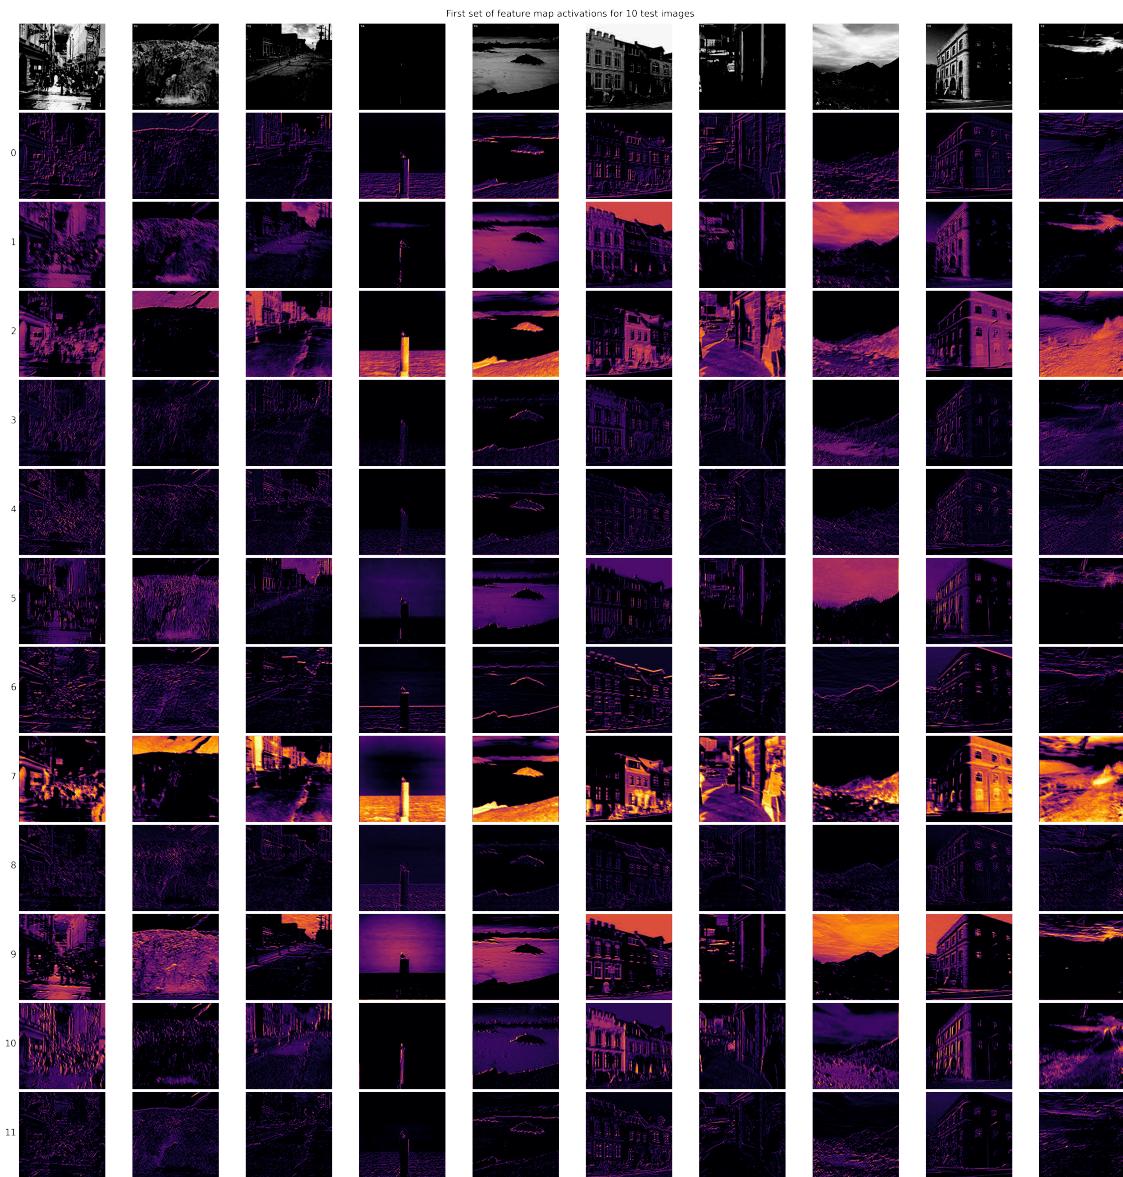
    # show the original picture
    img = X[pici, 0, :, :].detach().cpu()
    axs[0, pici].imshow(img, cmap="gray", vmin=0, vmax=1)
    axs[0, pici].axis('off')
    axs[0, pici].text(2,2,'T:%s'%int(y[pici].
    ↪item()),ha='left',va='top',color='w',fontweight='bold')

    for feati in range(feature_maps1.shape[1]):
        # extract the feature map from this image
        img = feature_maps1[pici,feati,:,:].detach().cpu()
        axs[feati+1,pici].imshow(img,cmap='inferno',vmin=0,vmax=torch.max(img)*.
        ↪9)
        axs[feati+1,pici].axis('off')
        axs[feati+1,pici].text(-5,feature_maps1.shape[2]/2,feati,ha='right',□
        ↪fontsize=40) if pici==0 else None

plt.tight_layout()
plt.suptitle(f'First set of feature map activations for {number_of_images} test□
    ↪images',x=.5,y=1.01, fontsize=40)

# save the plot
# plt.savefig('feature_maps1.png', dpi=300, bbox_inches='tight')
plt.show()

```



```
[18]: # Feature maps from the convolution layer 2
```

```
number_of_images = 10

fig,axs = plt.subplots(feature_maps2.shape[1] + 1, number_of_images,□
    figsize=(70, 70))

for pici in range(number_of_images):

    # show the original picture
    img = X[pici, 0, :, :].detach().cpu()
```

```

axs[0, pici].imshow(img, cmap="gray", vmin=0, vmax=1)
axs[0, pici].axis('off')
axs[0, pici].text(2,2,'T:%s'%int(y[pici].
↪item()),ha='left',va='top',color='w',fontweight='bold')

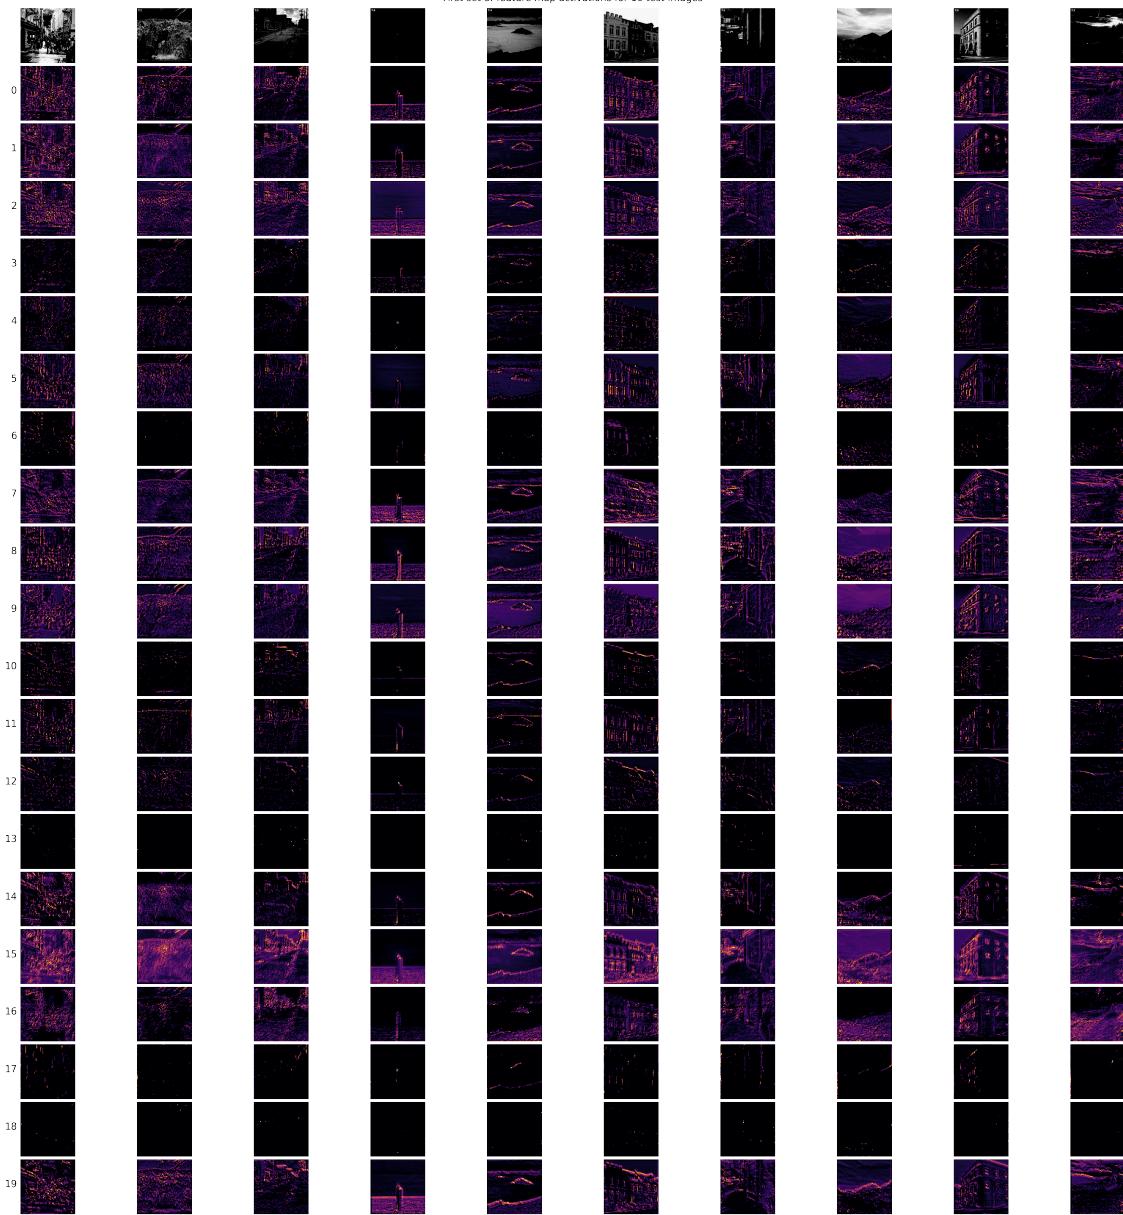
for feati in range(feature_maps2.shape[1]):
    # extract the feature map from this image
    img = feature_maps2[pici,feati,:,:].detach().cpu()
    axs[feati+1,pici].imshow(img,cmap='inferno',vmin=0,vmax=torch.max(img)*.
↪9)
    axs[feati+1,pici].axis('off')
    axs[feati+1,pici].text(-5,feature_maps2.shape[2]/2,feati,ha='right',u
↪fontsize=40) if pici==0 else None

plt.tight_layout()
plt.suptitle(f'First set of feature map activations for {number_of_images} test_
↪images',x=.5,y=1.01, fontsize=40)

# save the plot
# plt.savefig('feature_maps2.png', dpi=300, bbox_inches='tight')
plt.show()

```

First set of feature map activations for 10 test images



[]: