

March 5, 2023

1 NAME- BACHU GANESH

2 ROLL NUMBER-CH.EN.U4AIE20003

Data pipeline

Baseline model: train a simple CNN from scratch
 Transfer learning: pretrainend ConvNet as a feature extractor
 Transfer learning: fine-tune a pretrained ConvNet
 Test accuracy & visualize predictions

```
[1]: # Enable TensorFlow 2.0
      #%tensorflow_version 2.x
```

Colab only includes TensorFlow 2.x; %tensorflow_version has no effect.

```
[1]: !pip install tensorflow-gpu==2.11
      # Import Tensorflow
      import tensorflow as tf
      tf.__version__
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>
 Requirement already satisfied: tensorflow-gpu==2.11 in
 /usr/local/lib/python3.8/dist-packages (2.11.0)
 Requirement already satisfied: setuptools in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (57.4.0)
 Requirement already satisfied: numpy>=1.20 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (1.22.4)
 Requirement already satisfied: gast<=0.4.0,>=0.2.1 in
 /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (0.4.0)
 Requirement already satisfied: grpcio<2.0,>=1.24.3 in
 /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (1.51.3)
 Requirement already satisfied: keras<2.12,>=2.11.0 in
 /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (2.11.0)
 Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (1.4.0)
 Requirement already satisfied: opt-einsum>=2.3.2 in
 /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (3.3.0)
 Requirement already satisfied: google-pasta>=0.1.1 in

/usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (0.2.0)
 Requirement already satisfied: h5py>=2.9.0 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (3.1.0)
 Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (1.15.0)
 Requirement already satisfied: tensorflow-estimator<2.12,>=2.11.0 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (2.11.0)
 Requirement already satisfied: flatbuffers>=2.0 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (23.1.21)
 Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (1.15.0)
 Requirement already satisfied: protobuf<3.20,>=3.9.2 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (3.19.6)
 Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (0.31.0)
 Requirement already satisfied: tensorboard<2.12,>=2.11 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (2.11.2)
 Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (2.2.0)
 Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (1.6.3)
 Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (4.5.0)
 Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (15.0.6.1)
 Requirement already satisfied: packaging in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (23.0)
 Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.8/dist-packages (from astunparse>=1.6.0->tensorflow-gpu==2.11) (0.38.4)
 Requirement already satisfied: google-auth<3,>=1.6.3 in /usr/local/lib/python3.8/dist-packages (from tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (2.16.1)
 Requirement already satisfied: tensorboard-data-server<0.7.0,>=0.6.0 in /usr/local/lib/python3.8/dist-packages (from tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (0.6.1)
 Requirement already satisfied: tensorboard-plugin-wit>=1.6.0 in /usr/local/lib/python3.8/dist-packages (from tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (1.8.1)
 Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.8/dist-packages (from tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (2.2.3)
 Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.8/dist-packages (from tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (3.4.1)
 Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.8/dist-packages (from tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (2.25.1)
 Requirement already satisfied: google-auth-oauthlib<0.5,>=0.4.1 in /usr/local/lib/python3.8/dist-packages (from

tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (0.4.6)
 Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.8/dist-packages (from google-auth<3,>=1.6.3->tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (4.9)
 Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.8/dist-packages (from google-auth<3,>=1.6.3->tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (0.2.8)
 Requirement already satisfied: cachetools<6.0,>=2.0.0 in /usr/local/lib/python3.8/dist-packages (from google-auth<3,>=1.6.3->tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (5.3.0)
 Requirement already satisfied: requests-oauthlib>=0.7.0 in /usr/local/lib/python3.8/dist-packages (from google-auth-oauthlib<0.5,>=0.4.1->tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (1.3.1)
 Requirement already satisfied: importlib-metadata>=4.4 in /usr/local/lib/python3.8/dist-packages (from markdown>=2.6.8->tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (6.0.0)
 Requirement already satisfied: chardet<5,>=3.0.2 in /usr/local/lib/python3.8/dist-packages (from requests<3,>=2.21.0->tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (4.0.0)
 Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.8/dist-packages (from requests<3,>=2.21.0->tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (1.26.14)
 Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.8/dist-packages (from requests<3,>=2.21.0->tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (2022.12.7)
 Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.8/dist-packages (from requests<3,>=2.21.0->tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (2.10)
 Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.8/dist-packages (from werkzeug>=1.0.1->tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (2.1.2)
 Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.8/dist-packages (from importlib-metadata>=4.4->markdown>=2.6.8->tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (3.15.0)
 Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in /usr/local/lib/python3.8/dist-packages (from pyasn1-modules>=0.2.1->google-auth<3,>=1.6.3->tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (0.4.8)
 Requirement already satisfied: oauthlib>=3.0.0 in /usr/local/lib/python3.8/dist-packages (from requests-oauthlib>=0.7.0->google-auth-oauthlib<0.5,>=0.4.1->tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (3.2.2)

[1]: '2.11.0'

```

[2]: # Import TensorFlow datasets
import tensorflow_datasets as tfds
tfds.disable_progress_bar()
  
```

```

# Import Keras
from tensorflow import keras
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv2D, Flatten, Dropout,
    ↳MaxPooling2D, GlobalAveragePooling2D
from tensorflow.keras.optimizers import Adam

# Import Numpy
import numpy as np
import matplotlib.pyplot as plt

```

3 Data pipeline

Load the tf_flowers dataset tf_flowers is one of the TensorFlow 2.0 datasets with 3670 samples.

```

[3]: # Load train and validation datasets
(raw_train, raw_validation, raw_test), metadata = tfds.load(
    name='tf_flowers',
    split=['train[:80%]', 'train[80%:90%]', 'train[90%:]'], # train,
    ↳validation, test split of 8:1:1
    with_info=True,
    as_supervised=True)

```

Downloading and preparing dataset Unknown size (download: Unknown size, generated: Unknown size, total: Unknown size) to /root/tensorflow_datasets/tf_flowers/3.0.1... Dataset tf_flowers downloaded and prepared to /root/tensorflow_datasets/tf_flowers/3.0.1. Subsequent calls will reuse this data.

```

[4]: print("Total number of samples:", metadata.splits['train'].num_examples)

```

Total number of samples: 3670

```

[5]: num_classes = metadata.features['label'].num_classes
num_train = len(list(raw_train))
num_validation = len(list(raw_validation))
num_test = len(list(raw_test))

print("Number of classes:", num_classes)
print("Number of training samples:", num_train)
print("Number of validation samples:", num_validation)
print("Number of test samples:", num_test)

```

Number of classes: 5

Number of training samples: 2936
Number of validation samples: 367
Number of test samples: 367

```
[6]: # Inspect datasets before data preprocessing
print(raw_train)
print(raw_validation)
print(raw_test)
```

```
<PrefetchDataset element_spec=(TensorSpec(shape=(None, None, 3), dtype=tf.uint8,
name=None), TensorSpec(shape=(), dtype=tf.int64, name=None))>
<PrefetchDataset element_spec=(TensorSpec(shape=(None, None, 3), dtype=tf.uint8,
name=None), TensorSpec(shape=(), dtype=tf.int64, name=None))>
<PrefetchDataset element_spec=(TensorSpec(shape=(None, None, 3), dtype=tf.uint8,
name=None), TensorSpec(shape=(), dtype=tf.int64, name=None))>
```

```
[7]: # Get labels / class names
class_names = np.array(metadata.features['label'].names)
print(class_names)
```


```
['dandelion' 'daisy' 'tulips' 'sunflowers' 'roses']
```

4 Visualize the data

Let's take a look a few of the flower images

```
[8]: label_names = metadata.features['label'].int2str

for image, label in raw_train.take(2):
    plt.figure()
    plt.imshow(image)
    plt.title(label_names[label])
```



5 Image preprocessing

Resize, normalize, augment, shuffle and batch the data.

Resize and normalize dataset

```
[9]: IMG_SIZE = 224
      IMG_SHAPE = (IMG_SIZE, IMG_SIZE, 3)

      def format_example(image, label):
          image = tf.cast(image, tf.float32)
          image = tf.image.resize(image, (IMG_SIZE, IMG_SIZE))
          image = image/255.0
          return image, label
```

```
[10]: train = raw_train.map(format_example)
      validation = raw_validation.map(format_example)
      test = raw_test.map(format_example)
```

```
[11]: def augment_data(image, label):
      image = tf.image.random_flip_left_right(image)
      image = tf.image.random_contrast(image, lower=0.0, upper=1.0)
      image = tf.stack(image, axis=0)
      image = tf.image.random_crop(image, size=[IMG_SIZE, IMG_SIZE, 3])
      return image, label
```

```
[12]: train = train.map(augment_data)
```

Shuffle and batch dataset

```
[13]: BATCH_SIZE = 32
      SHUFFLE_BUFFER_SIZE = 1000

      train_batches = train.shuffle(SHUFFLE_BUFFER_SIZE).batch(BATCH_SIZE).repeat()
      validation_batches = validation.batch(BATCH_SIZE).repeat()
      test_batches = test.batch(BATCH_SIZE)
```

```
[14]: # Inspect datasets after data preprocessing
      print(train_batches)
      print(validation_batches)
      print(test_batches)
```

```
<RepeatDataset element_spec=(TensorSpec(shape=(None, 224, 224, 3),
dtype=tf.float32, name=None), TensorSpec(shape=(None,), dtype=tf.int64,
name=None))>
<RepeatDataset element_spec=(TensorSpec(shape=(None, 224, 224, 3),
dtype=tf.float32, name=None), TensorSpec(shape=(None,), dtype=tf.int64,
name=None))>
<BatchDataset element_spec=(TensorSpec(shape=(None, 224, 224, 3),
```

```
dtype=tf.float32, name=None), TensorSpec(shape=(None,), dtype=tf.int64,
name=None))>
```

```
[15]: # Inspect a batch of data
for image_batch, label_batch in train_batches.take(1):
    pass

image_batch.shape
```

```
[15]: TensorShape([32, 224, 224, 3])
```

6 Training

```
[16]: # Set training parameters
NUM_EPOCHS = 10
steps_per_epoch = round(num_train)//BATCH_SIZE
validation_steps = round(num_validation)//BATCH_SIZE
```

```
[18]: # Display training curves
def display_training_curves(history, title):
    acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']

    loss = history.history['loss']
    val_loss = history.history['val_loss']

    epochs_range = range(NUM_EPOCHS)

    plt.plot(epochs_range, acc, label='Train accuracy')
    plt.plot(epochs_range, val_acc, label='Val accuracy')
    plt.title(title)
    plt.legend(loc='upper left')
    plt.figure()

    plt.show()
```

7 Baseline - train from scratch

Train a very simple CNN model and use the accuracy metrics as baseline to compare with transfer learning results.

Create model

```
[19]: def build_model_from_scratch():

    model = Sequential([
```



```

# Must define the input shape in the first layer of the neural network
Conv2D(filters=32, kernel_size=3, padding='same', activation='relu',
↪input_shape=IMG_SHAPE),
MaxPooling2D(pool_size=2),

Conv2D(filters=64, kernel_size=3, padding='same', activation='relu'),
MaxPooling2D(pool_size=2),

Flatten(),
Dense(64, activation='relu'),
Dense(num_classes, activation='softmax')
])

return model

```

```
[20]: simple_cnn_model = build_model_from_scratch()
```

```
[21]: simple_cnn_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 224, 224, 32)	896
max_pooling2d (MaxPooling2D)	(None, 112, 112, 32)	0
conv2d_1 (Conv2D)	(None, 112, 112, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 64)	0
flatten (Flatten)	(None, 200704)	0
dense (Dense)	(None, 64)	12845120
dense_1 (Dense)	(None, 5)	325

=====
 Total params: 12,864,837
 Trainable params: 12,864,837
 Non-trainable params: 0
 =====

Compile and train the model

```
[22]: def train_model(model):
    model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

    history = model.fit(train_batches,
                        epochs=NUM_EPOCHS,
                        validation_data=validation_batches,
                        steps_per_epoch=steps_per_epoch,
                        validation_steps=validation_steps)


    return history
```

```
[23]: %%time
history = train_model(simple_cnn_model)
```

```
Epoch 1/10
91/91 [=====] - 17s 58ms/step - loss: 2.0097 -
accuracy: 0.3527 - val_loss: 1.2163 - val_accuracy: 0.5142
Epoch 2/10
91/91 [=====] - 7s 71ms/step - loss: 1.3237 - accuracy:
0.4415 - val_loss: 1.1904 - val_accuracy: 0.5398
Epoch 3/10
91/91 [=====] - 5s 59ms/step - loss: 1.2410 - accuracy:
0.4917 - val_loss: 1.2132 - val_accuracy: 0.5568
Epoch 4/10
91/91 [=====] - 6s 62ms/step - loss: 1.1657 - accuracy:
0.5234 - val_loss: 1.4410 - val_accuracy: 0.4773
Epoch 5/10
91/91 [=====] - 6s 71ms/step - loss: 1.1207 - accuracy:
0.5386 - val_loss: 1.1207 - val_accuracy: 0.5938
Epoch 6/10
91/91 [=====] - 6s 62ms/step - loss: 1.0637 - accuracy:
0.5740 - val_loss: 1.0783 - val_accuracy: 0.6222
Epoch 7/10
91/91 [=====] - 7s 76ms/step - loss: 0.9996 - accuracy:
0.5992 - val_loss: 1.2210 - val_accuracy: 0.6222
Epoch 8/10
91/91 [=====] - 5s 54ms/step - loss: 0.9551 - accuracy:
0.6240 - val_loss: 1.2361 - val_accuracy: 0.6506
Epoch 9/10
91/91 [=====] - 5s 53ms/step - loss: 0.8441 - accuracy:
0.6742 - val_loss: 1.1498 - val_accuracy: 0.6449
Epoch 10/10
91/91 [=====] - 7s 72ms/step - loss: 0.8217 - accuracy:
0.6983 - val_loss: 1.1498 - val_accuracy: 0.6534
CPU times: user 1min 23s, sys: 4.6 s, total: 1min 28s
```

Wall time: 1min 22s

```
[24]: # Display training curve
display_training_curves(history, "A Simple CNN Model")
```



<Figure size 432x288 with 0 Axes>

8 Transfer learning

Now let's see how transfer learning can help achieve better results.

Feature extractor Use MobileNetV2 as a feature extractor and add a classifier on top of it.

Create base model

```
[25]: # Create base model from tf.keras pre-trained model MobileNetV2
base_model = MobileNetV2(input_shape=IMG_SHAPE,
                          weights="imagenet",
                          include_top=False)
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_224_no_top.h5

9406464/9406464 [=====] - 0s 0us/step

Freeze all layers of the base model

```
[31]: base_model.trainable = False
```

Add a classifier head

Create a new model by adding a classifier on top of the base model.

```
[26]: def build_mobilenetv2_model(base_model):
    model = Sequential([
        base_model,
        Conv2D(32, 3, activation='relu'),
        GlobalAveragePooling2D(),
        Dense(num_classes, activation='softmax')]
    )

    return model
```

```
[27]: model = build_mobilenetv2_model(base_model)
```

Compile the model

```
[28]: model.compile(optimizer= Adam(),
                    loss='sparse_categorical_crossentropy',
                    metrics=['accuracy'])
```

```
[29]: model.summary()
```

Model: "sequential_1"


Layer (type)	Output Shape	Param #
mobilenetv2_1.00_224 (Functional)	(None, 7, 7, 1280)	2257984
conv2d_2 (Conv2D)	(None, 5, 5, 32)	368672
global_average_pooling2d (GlobalAveragePooling2D)	(None, 32)	0
dense_2 (Dense)	(None, 5)	165
Total params: 2,626,821		
Trainable params: 2,592,709		
Non-trainable params: 34,112		

Train the model

```
[34]: %%time
      history = model.fit(train_batches,
                          epochs=NUM_EPOCHS,
                          validation_data=validation_batches,
                          steps_per_epoch=steps_per_epoch,
                          validation_steps=validation_steps)

Epoch 1/10
91/91 [=====] - 14s 136ms/step - loss: 0.2709 -
accuracy: 0.9076 - val_loss: 8.4277 - val_accuracy: 0.3608
Epoch 2/10
91/91 [=====] - 14s 151ms/step - loss: 0.2084 -
accuracy: 0.9298 - val_loss: 3.0599 - val_accuracy: 0.5852
Epoch 3/10
91/91 [=====] - 14s 158ms/step - loss: 0.2512 -
accuracy: 0.9167 - val_loss: 10.5396 - val_accuracy: 0.2614
Epoch 4/10
91/91 [=====] - 14s 149ms/step - loss: 0.1741 -
accuracy: 0.9421 - val_loss: 3.3062 - val_accuracy: 0.5426
Epoch 5/10
91/91 [=====] - 13s 146ms/step - loss: 0.1817 -
accuracy: 0.9377 - val_loss: 4.2440 - val_accuracy: 0.4602
Epoch 6/10
91/91 [=====] - 13s 143ms/step - loss: 0.1781 -
accuracy: 0.9421 - val_loss: 4.0010 - val_accuracy: 0.5881
Epoch 7/10
91/91 [=====] - 13s 142ms/step - loss: 0.1924 -
accuracy: 0.9342 - val_loss: 4.7941 - val_accuracy: 0.5284
Epoch 8/10
91/91 [=====] - 13s 140ms/step - loss: 0.2072 -
accuracy: 0.9308 - val_loss: 2.9679 - val_accuracy: 0.5568
Epoch 9/10
91/91 [=====] - 13s 139ms/step - loss: 0.1633 -
accuracy: 0.9387 - val_loss: 5.6628 - val_accuracy: 0.4432
Epoch 10/10
91/91 [=====] - 14s 149ms/step - loss: 0.1506 -
accuracy: 0.9456 - val_loss: 3.4156 - val_accuracy: 0.7102
CPU times: user 2min 28s, sys: 4.96 s, total: 2min 33s
Wall time: 2min 21s
```

```
[35]: # Display training curve
      display_training_curves(history, "Pretrained ConvNet as Feature Extractor")
```



<Figure size 432x288 with 0 Axes>

[]: