

March 5, 2023

## 1 NAME- BACHU GANESH

## 2 ROLL NUMBER-CH.EN.U4AIE20003

### Data pipeline

Baseline model: train a simple CNN from scratch  
Transfer learning: pretrainend ConvNet as a feature extractor  
Transfer learning: fine-tune a pretrained ConvNet  
Test accuracy & visualize predictions

```
[1]: # Enable TensorFlow 2.0
      #%tensorflow_version 2.x
```

Colab only includes TensorFlow 2.x; %tensorflow\_version has no effect.

```
[2]: !pip install tensorflow-gpu==2.11
      # Import Tensorflow
      import tensorflow as tf
      tf.__version__
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Collecting tensorflow-gpu==2.11

Downloading

tensorflow\_gpu-2.11.0-cp38-cp38-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl  
(588.3 MB)

588.3/588.3

MB 2.9 MB/s eta 0:00:00

Requirement already satisfied: google-pasta>=0.1.1 in

/usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (0.2.0)

Requirement already satisfied: protobuf<3.20,>=3.9.2 in

/usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (3.19.6)

Requirement already satisfied: flatbuffers>=2.0 in

/usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (23.1.21)

Requirement already satisfied: libclang>=13.0.0 in

/usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (15.0.6.1)

Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in

/usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (0.31.0)

Requirement already satisfied: h5py>=2.9.0 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (3.1.0)

Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (1.15.0)

Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (2.2.0)

Requirement already satisfied: setuptools in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (57.4.0)

Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (1.6.3)

Requirement already satisfied: packaging in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (23.0)

Requirement already satisfied: tensorflow-estimator<2.12,>=2.11.0 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (2.11.0)

Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (3.3.0)

Requirement already satisfied: numpy>=1.20 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (1.22.4)

Requirement already satisfied: tensorboard<2.12,>=2.11 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (2.11.2)

Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (4.5.0)

Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (1.51.3)

Requirement already satisfied: gast<=0.4.0,>=0.2.1 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (0.4.0)

Requirement already satisfied: keras<2.12,>=2.11.0 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (2.11.0)

Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (1.4.0)

Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (1.15.0)

Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.8/dist-packages (from astunparse>=1.6.0->tensorflow-gpu==2.11) (0.38.4)

Requirement already satisfied: google-auth-oauthlib<0.5,>=0.4.1 in /usr/local/lib/python3.8/dist-packages (from tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (0.4.6)

Requirement already satisfied: google-auth<3,>=1.6.3 in /usr/local/lib/python3.8/dist-packages (from tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (2.16.1)

Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.8/dist-packages (from tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (2.25.1)

Requirement already satisfied: tensorboard-plugin-wit>=1.6.0 in /usr/local/lib/python3.8/dist-packages (from tensorboard<2.12,>=2.11->tensorflow-gpu==2.11) (1.8.1)

Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.8/dist-

packages (from tensorflow-gpu==2.11) (2.2.3)  
 Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (3.4.1)  
 Requirement already satisfied: tensorboard-data-server<0.7.0,>=0.6.0 in /usr/local/lib/python3.8/dist-packages (from tensorflow-gpu==2.11) (0.6.1)  
 Requirement already satisfied: cachetools<6.0,>=2.0.0 in /usr/local/lib/python3.8/dist-packages (from google-auth<3,>=1.6.3->tensorflow-gpu==2.11) (5.3.0)  
 Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.8/dist-packages (from google-auth<3,>=1.6.3->tensorflow-gpu==2.11) (0.2.8)  
 Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.8/dist-packages (from google-auth<3,>=1.6.3->tensorflow-gpu==2.11) (4.9)  
 Requirement already satisfied: requests-oauthlib>=0.7.0 in /usr/local/lib/python3.8/dist-packages (from google-auth-oauthlib<0.5,>=0.4.1->tensorflow-gpu==2.11) (1.3.1)  
 Requirement already satisfied: importlib-metadata>=4.4 in /usr/local/lib/python3.8/dist-packages (from markdown>=2.6.8->tensorflow-gpu==2.11) (6.0.0)  
 Requirement already satisfied: chardet<5,>=3.0.2 in /usr/local/lib/python3.8/dist-packages (from requests<3,>=2.21.0->tensorflow-gpu==2.11) (4.0.0)  
 Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.8/dist-packages (from requests<3,>=2.21.0->tensorflow-gpu==2.11) (2.10)  
 Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.8/dist-packages (from requests<3,>=2.21.0->tensorflow-gpu==2.11) (1.26.14)  
 Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.8/dist-packages (from requests<3,>=2.21.0->tensorflow-gpu==2.11) (2022.12.7)  
 Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.8/dist-packages (from werkzeug>=1.0.1->tensorflow-gpu==2.11) (2.1.2)  
 Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.8/dist-packages (from importlib-metadata>=4.4->markdown>=2.6.8->tensorflow-gpu==2.11) (3.15.0)  
 Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in /usr/local/lib/python3.8/dist-packages (from pyasn1-modules>=0.2.1->google-auth<3,>=1.6.3->tensorflow-gpu==2.11) (0.4.8)  
 Requirement already satisfied: oauthlib>=3.0.0 in /usr/local/lib/python3.8/dist-packages (from requests-oauthlib>=0.7.0->google-auth-oauthlib<0.5,>=0.4.1->tensorflow-gpu==2.11) (3.2.2)  
 Installing collected packages: tensorflow-gpu  
 Successfully installed tensorflow-gpu-2.11.0

```
[2]: '2.11.0'
```

```
[3]: # Import TensorFlow datasets
import tensorflow_datasets as tfds
tfds.disable_progress_bar()

# Import Keras
from tensorflow import keras
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv2D, Flatten, Dropout,
↳MaxPooling2D, GlobalAveragePooling2D
from tensorflow.keras.optimizers import Adam

# Import Numpy
import numpy as np
import matplotlib.pyplot as plt
```

### 3 Data pipeline

Load the tf\_flowers dataset tf\_flowers is one of the TensorFlow 2.0 datasets with 3670 samples.

```
[4]: # Load train and validation datasets
(raw_train, raw_validation, raw_test), metadata = tfds.load(
    name='tf_flowers',
    split=['train[:80%]', 'train[80%:90%]', 'train[90%:]'], # train,
↳validation, test split of 8:1:1
    with_info=True,
    as_supervised=True)
```

Downloading and preparing dataset Unknown size (download: Unknown size, generated: Unknown size, total: Unknown size) to /root/tensorflow\_datasets/tf\_flowers/3.0.1...  
Dataset tf\_flowers downloaded and prepared to /root/tensorflow\_datasets/tf\_flowers/3.0.1. Subsequent calls will reuse this data.

```
[5]: print("Total number of samples:", metadata.splits['train'].num_examples)
```

Total number of samples: 3670

```
[6]: num_classes = metadata.features['label'].num_classes
num_train = len(list(raw_train))
num_validation = len(list(raw_validation))
num_test = len(list(raw_test))

print("Number of classes:", num_classes)
```

```
print("Number of training samples:", num_train)
print("Number of validation samples:", num_validation)
print("Number of test samples:", num_test)
```

```
Number of classes: 5
Number of training samples: 2936
Number of validation samples: 367
Number of test samples: 367
```

```
[7]: # Inspect datasets before data preprocessing
print(raw_train)
print(raw_validation)
print(raw_test)
```

```
<PrefetchDataset element_spec=(TensorSpec(shape=(None, None, 3), dtype=tf.uint8,
name=None), TensorSpec(shape=(), dtype=tf.int64, name=None))>
<PrefetchDataset element_spec=(TensorSpec(shape=(None, None, 3), dtype=tf.uint8,
name=None), TensorSpec(shape=(), dtype=tf.int64, name=None))>
<PrefetchDataset element_spec=(TensorSpec(shape=(None, None, 3), dtype=tf.uint8,
name=None), TensorSpec(shape=(), dtype=tf.int64, name=None))>
```

```
[8]: # Get labels / class names
class_names = np.array(metadata.features['label'].names)
print(class_names)
```


```
['dandelion' 'daisy' 'tulips' 'sunflowers' 'roses']
```

## 4 Visualize the data

Let's take a look a few of the flower images

```
[9]: label_names = metadata.features['label'].int2str

for image, label in raw_train.take(2):
    plt.figure()
    plt.imshow(image)
    plt.title(label_names(label))
```



## 5 Image preprocessing

Resize, normalize, augment, shuffle and batch the data.

Resize and normalize dataset

```
[10]: IMG_SIZE = 224
      IMG_SHAPE = (IMG_SIZE, IMG_SIZE, 3)

      def format_example(image, label):
          image = tf.cast(image, tf.float32)
          image = tf.image.resize(image, (IMG_SIZE, IMG_SIZE))
          image = image/255.0
          return image, label
```

```
[11]: train = raw_train.map(format_example)
      validation = raw_validation.map(format_example)
      test = raw_test.map(format_example)
```

```
[12]: def augment_data(image, label):
      image = tf.image.random_flip_left_right(image)
      image = tf.image.random_contrast(image, lower=0.0, upper=1.0)
      image = tf.stack(image, axis=0)
      image = tf.image.random_crop(image, size=[IMG_SIZE, IMG_SIZE, 3])
      return image, label
```

```
[13]: train = train.map(augment_data)
```

Shuffle and batch dataset

```
[14]: BATCH_SIZE = 32
      SHUFFLE_BUFFER_SIZE = 1000

      train_batches = train.shuffle(SHUFFLE_BUFFER_SIZE).batch(BATCH_SIZE).repeat()
      validation_batches = validation.batch(BATCH_SIZE).repeat()
      test_batches = test.batch(BATCH_SIZE)
```

```
[15]: # Inspect datasets after data preprocessing
      print(train_batches)
      print(validation_batches)
      print(test_batches)
```

```
<RepeatDataset element_spec=(TensorSpec(shape=(None, 224, 224, 3),
dtype=tf.float32, name=None), TensorSpec(shape=(None,), dtype=tf.int64,
name=None))>
<RepeatDataset element_spec=(TensorSpec(shape=(None, 224, 224, 3),
dtype=tf.float32, name=None), TensorSpec(shape=(None,), dtype=tf.int64,
name=None))>
<BatchDataset element_spec=(TensorSpec(shape=(None, 224, 224, 3),
```

```
dtype=tf.float32, name=None), TensorSpec(shape=(None,), dtype=tf.int64,
name=None))>
```

```
[16]: # Inspect a batch of data
for image_batch, label_batch in train_batches.take(1):
    pass

image_batch.shape
```

```
[16]: TensorShape([32, 224, 224, 3])
```

## 6 Training

```
[17]: # Set training parameters
NUM_EPOCHS = 10
steps_per_epoch = round(num_train)//BATCH_SIZE
validation_steps = round(num_validation)//BATCH_SIZE
```

```
[18]: # Display training curves
def display_training_curves(history, title):
    acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']

    loss = history.history['loss']
    val_loss = history.history['val_loss']

    epochs_range = range(NUM_EPOCHS)

    plt.plot(epochs_range, acc, label='Train accuracy')
    plt.plot(epochs_range, val_acc, label='Val accuracy')
    plt.title(title)
    plt.legend(loc='upper left')
    plt.figure()

    plt.show()
```

## 7 Baseline - train from scratch

Train a very simple CNN model and use the accuracy metrics as baseline to compare with transfer learning results.

### Create model

```
[19]: def build_model_from_scratch():

    model = Sequential([
```



```

# Must define the input shape in the first layer of the neural network
Conv2D(filters=32, kernel_size=3, padding='same', activation='relu',
↪input_shape=IMG_SHAPE),
MaxPooling2D(pool_size=2),

Conv2D(filters=64, kernel_size=3, padding='same', activation='relu'),
MaxPooling2D(pool_size=2),

Flatten(),
Dense(64, activation='relu'),
Dense(num_classes, activation='softmax')
])

return model

```

```
[20]: simple_cnn_model = build_model_from_scratch()
```

```
[21]: simple_cnn_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 224, 224, 32)	896
max_pooling2d (MaxPooling2D)	(None, 112, 112, 32)	0
conv2d_1 (Conv2D)	(None, 112, 112, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 64)	0
flatten (Flatten)	(None, 200704)	0
dense (Dense)	(None, 64)	12845120
dense_1 (Dense)	(None, 5)	325
Total params: 12,864,837		
Trainable params: 12,864,837		
Non-trainable params: 0		

Compile and train the model

```
[22]: def train_model(model):
    model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

    history = model.fit(train_batches,
                        epochs=NUM_EPOCHS,
                        validation_data=validation_batches,
                        steps_per_epoch=steps_per_epoch,
                        validation_steps=validation_steps)


    return history
```

```
[23]: %%time
history = train_model(simple_cnn_model)
```

```
Epoch 1/10
91/91 [=====] - 16s 58ms/step - loss: 1.9384 -
accuracy: 0.3190 - val_loss: 1.3401 - val_accuracy: 0.4261
Epoch 2/10
91/91 [=====] - 7s 70ms/step - loss: 1.2938 - accuracy:
0.4483 - val_loss: 1.1256 - val_accuracy: 0.5739
Epoch 3/10
91/91 [=====] - 5s 56ms/step - loss: 1.2086 - accuracy:
0.5048 - val_loss: 1.0393 - val_accuracy: 0.5767
Epoch 4/10
91/91 [=====] - 6s 65ms/step - loss: 1.1559 - accuracy:
0.5341 - val_loss: 1.1216 - val_accuracy: 0.5540
Epoch 5/10
91/91 [=====] - 6s 66ms/step - loss: 1.1041 - accuracy:
0.5475 - val_loss: 1.2216 - val_accuracy: 0.5398
Epoch 6/10
91/91 [=====] - 5s 53ms/step - loss: 1.0434 - accuracy:
0.5851 - val_loss: 1.0281 - val_accuracy: 0.6364
Epoch 7/10
91/91 [=====] - 5s 58ms/step - loss: 0.9753 - accuracy:
0.6119 - val_loss: 0.9794 - val_accuracy: 0.6477
Epoch 8/10
91/91 [=====] - 5s 59ms/step - loss: 0.9097 - accuracy:
0.6481 - val_loss: 1.3184 - val_accuracy: 0.5653
Epoch 9/10
91/91 [=====] - 5s 52ms/step - loss: 0.8222 - accuracy:
0.6808 - val_loss: 1.1413 - val_accuracy: 0.6392
Epoch 10/10
91/91 [=====] - 6s 64ms/step - loss: 0.7634 - accuracy:
0.7097 - val_loss: 1.2268 - val_accuracy: 0.6222
CPU times: user 1min 23s, sys: 5.52 s, total: 1min 28s
```

Wall time: 1min 5s

```
[24]: # Display training curve
display_training_curves(history, "A Simple CNN Model")
```



<Figure size 432x288 with 0 Axes>

## 8 Transfer learning

Now let's see how transfer learning can help achieve better results.

**Feature extractor** Use MobileNetV2 as a feature extractor and add a classifier on top of it.

**Create base model**

```
[25]: # Create base model from tf.keras pre-trained model MobileNetV2
base_model = MobileNetV2(input_shape=IMG_SHAPE,
                          weights="imagenet",
                          include_top=False)
```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/mobilenet\\_v2/mobilenet\\_v2\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_1.0\\_224\\_no\\_top.h5](https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_224_no_top.h5)

9406464/9406464 [=====] - 0s 0us/step

**Freeze all layers of the base model**

```
[26]: base_model.trainable = False
```

### Add a classifier head

Create a new model by adding a classifier on top of the base model.

```
[27]: def build_mobilenetv2_model(base_model):
    model = Sequential([
        base_model,
        Conv2D(32, 3, activation='relu'),
        GlobalAveragePooling2D(),
        Dense(num_classes, activation='softmax')]
    )

    return model
```

```
[28]: model = build_mobilenetv2_model(base_model)
```

### Compile the model

```
[29]: model.compile(optimizer= Adam(),
                    loss='sparse_categorical_crossentropy',
                    metrics=['accuracy'])
```

```
[30]: model.summary()
```

Model: "sequential\_1"


Layer (type)	Output Shape	Param #
mobilenetv2_1.00_224 (Functional)	(None, 7, 7, 1280)	2257984
conv2d_2 (Conv2D)	(None, 5, 5, 32)	368672
global_average_pooling2d (GlobalAveragePooling2D)	(None, 32)	0
dense_2 (Dense)	(None, 5)	165
Total params: 2,626,821		
Trainable params: 368,837		
Non-trainable params: 2,257,984		

### Train the model

```
[31]: %%time
history = model.fit(train_batches,
                    epochs=NUM_EPOCHS,
                    validation_data=validation_batches,
                    steps_per_epoch=steps_per_epoch,
                    validation_steps=validation_steps)

Epoch 1/10
91/91 [=====] - 13s 86ms/step - loss: 0.9577 -
accuracy: 0.6075 - val_loss: 0.4419 - val_accuracy: 0.8608
Epoch 2/10
91/91 [=====] - 5s 53ms/step - loss: 0.5474 - accuracy:
0.7879 - val_loss: 0.3968 - val_accuracy: 0.8864
Epoch 3/10
91/91 [=====] - 6s 69ms/step - loss: 0.4731 - accuracy:
0.8271 - val_loss: 0.4475 - val_accuracy: 0.8722
Epoch 4/10
91/91 [=====] - 6s 69ms/step - loss: 0.4196 - accuracy:
0.8433 - val_loss: 0.3957 - val_accuracy: 0.9034
Epoch 5/10
91/91 [=====] - 6s 63ms/step - loss: 0.3663 - accuracy:
0.8598 - val_loss: 0.4629 - val_accuracy: 0.8722
Epoch 6/10
91/91 [=====] - 5s 58ms/step - loss: 0.3311 - accuracy:
0.8716 - val_loss: 0.4141 - val_accuracy: 0.9091
Epoch 7/10
91/91 [=====] - 5s 55ms/step - loss: 0.3174 - accuracy:
0.8815 - val_loss: 0.3956 - val_accuracy: 0.8835
Epoch 8/10
91/91 [=====] - 5s 54ms/step - loss: 0.2944 - accuracy:
0.8902 - val_loss: 0.4223 - val_accuracy: 0.8835
Epoch 9/10
91/91 [=====] - 6s 67ms/step - loss: 0.2925 - accuracy:
0.8829 - val_loss: 0.3939 - val_accuracy: 0.9006
Epoch 10/10
91/91 [=====] - 5s 52ms/step - loss: 0.2910 - accuracy:
0.8891 - val_loss: 0.4393 - val_accuracy: 0.8977
CPU times: user 1min 24s, sys: 5.66 s, total: 1min 29s
Wall time: 1min 2s
```

```
[32]: # Display training curve
display_training_curves(history, "Pretrained ConvNet as Feature Extractor")
```



<Figure size 432x288 with 0 Axes>

## 9 Fine tuning

Unfreeze top layers for fine tuning

```
[33]: # Unfreeze all layers in base model
base_model.trainable = True
```

```
[34]: # Let's take a look to see how many layers are in the base model
print("Number of layers in the base model: ", len(base_model.layers))
```

Number of layers in the base model: 154

```
[35]: # Fine-tune from this layer onwards
fine_tune_at = 100

# Freeze all the layers before the `fine_tune_at` layer
for layer in base_model.layers[:fine_tune_at]:
    layer.trainable = False
```

Compile the model

```
[36]: model.compile(loss='sparse_categorical_crossentropy',
                    optimizer = Adam(1e-5),
```

```
metrics=['accuracy'])
```

```
[37]: model.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
mobilenetv2_1.00_224 (Functional)	(None, 7, 7, 1280)	2257984
conv2d_2 (Conv2D)	(None, 5, 5, 32)	368672
global_average_pooling2d (GlobalAveragePooling2D)	(None, 32)	0
dense_2 (Dense)	(None, 5)	165

=====  
Total params: 2,626,821  
Trainable params: 2,230,277  
Non-trainable params: 396,544  
=====

-----  
Continue to train the model

```
[38]: INITIAL_EPOCHS = 10  
FINE_TUNE_EPOCHS = 10  
TOTAL_EPOCHS = INITIAL_EPOCHS + FINE_TUNE_EPOCHS #20
```

```
[39]: %%time  
history_fine = model.fit(train_batches,  
                          epochs=TOTAL_EPOCHS, #20  
                          initial_epoch=INITIAL_EPOCHS, #10  
                          validation_data=validation_batches,  
                          steps_per_epoch=steps_per_epoch,  
                          validation_steps=validation_steps)
```

```
Epoch 11/20
```

```
91/91 [=====] - 24s 102ms/step - loss: 0.5985 -  
accuracy: 0.7809 - val_loss: 0.4390 - val_accuracy: 0.8949
```

```
Epoch 12/20
```

```
91/91 [=====] - 6s 67ms/step - loss: 0.5067 - accuracy:  
0.8220 - val_loss: 0.4325 - val_accuracy: 0.8949
```

```
Epoch 13/20
```

```
91/91 [=====] - 8s 85ms/step - loss: 0.4269 - accuracy:  
0.8382 - val_loss: 0.4255 - val_accuracy: 0.8977
```

```
Epoch 14/20
```

```

91/91 [=====] - 7s 77ms/step - loss: 0.3753 - accuracy:
0.8640 - val_loss: 0.4247 - val_accuracy: 0.8977
Epoch 15/20
91/91 [=====] - 8s 83ms/step - loss: 0.3500 - accuracy:
0.8664 - val_loss: 0.4118 - val_accuracy: 0.9006
Epoch 16/20
91/91 [=====] - 7s 72ms/step - loss: 0.3290 - accuracy:
0.8778 - val_loss: 0.4066 - val_accuracy: 0.8977
Epoch 17/20
91/91 [=====] - 7s 79ms/step - loss: 0.3220 - accuracy:
0.8795 - val_loss: 0.3955 - val_accuracy: 0.9034
Epoch 18/20
91/91 [=====] - 7s 74ms/step - loss: 0.3008 - accuracy:
0.8888 - val_loss: 0.3937 - val_accuracy: 0.9034
Epoch 19/20
91/91 [=====] - 8s 84ms/step - loss: 0.3031 - accuracy:
0.8822 - val_loss: 0.3839 - val_accuracy: 0.9091
Epoch 20/20
91/91 [=====] - 8s 85ms/step - loss: 0.2931 - accuracy:
0.8881 - val_loss: 0.4000 - val_accuracy: 0.8977
CPU times: user 1min 45s, sys: 4.89 s, total: 1min 50s
Wall time: 1min 27s

```

```

[40]: # Display training curve

acc = history.history['accuracy'] + history_fine.history['accuracy']
val_acc = history.history['val_accuracy'] + history_fine.history['val_accuracy']

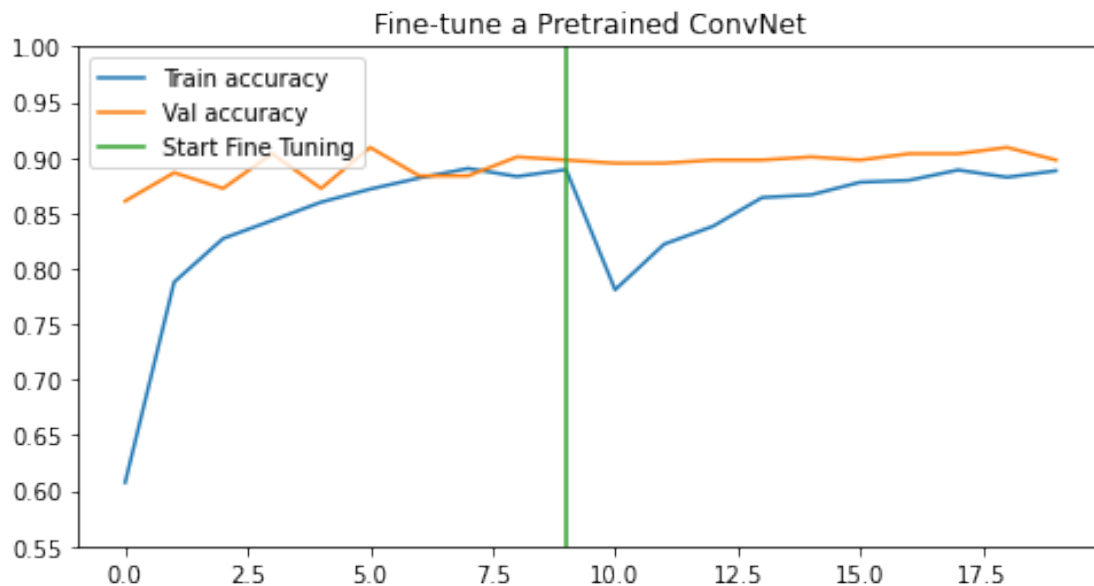
plt.figure(figsize=(8, 4))

plt.plot(acc, label='Train accuracy')
plt.plot(val_acc, label='Val accuracy')
plt.ylim([0.8, 1])
plt.plot([NUM_EPOCHS-1, NUM_EPOCHS-1], plt.ylim(ymin=0.55), label='Start Fine_
↳Tuning')
plt.title("Fine-tune a Pretrained ConvNet")
plt.legend(loc='upper left')

plt.show()

```





## 10 Test accuracy

```
[41]: # Evaluate the model on the test dataset
score = model.evaluate(test_batches, verbose=0)

# Print test accuracy
print('\n', 'Test accuracy:', score[1])
```

Test accuracy: 0.9046321511268616

## 11 Visualize predictions

First we get images and labels from a test batch, and then use the retrained model to make predictions.

```
[42]: image_batch, label_batch = next(iter(test_batches))

image_batch = image_batch.numpy()
label_batch = label_batch.numpy()

predicted_batch = model.predict(image_batch)
predicted_batch = tf.squeeze(predicted_batch).numpy()

predicted_class_ids = np.argmax(predicted_batch, axis=-1)
predicted_class_names = class_names[predicted_class_ids]
```

1/1 [=====] - 1s 1s/step

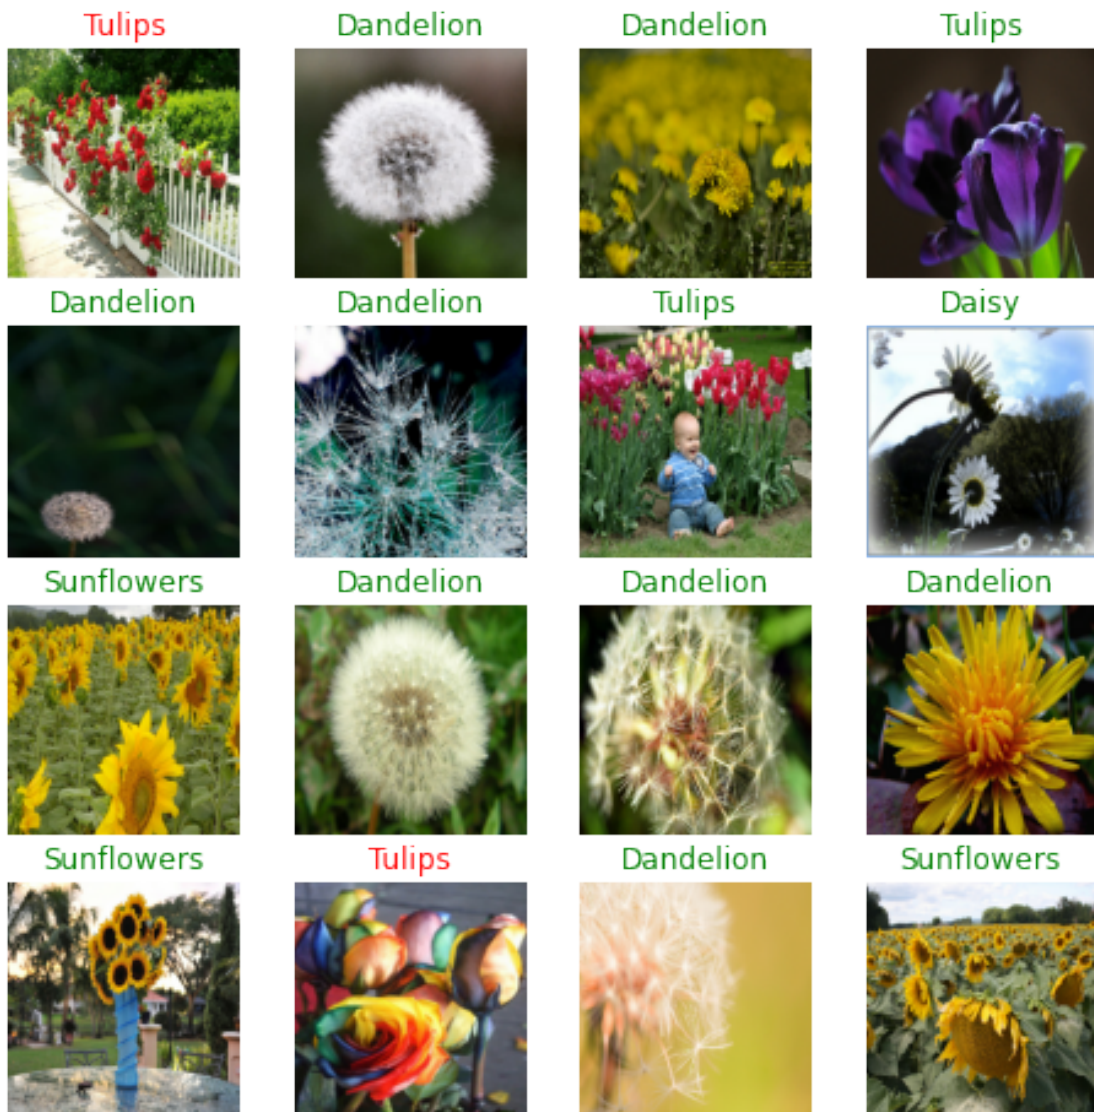
Then we visualize some of the images and compare whether the predicted labels equal to the true labels.

```
[43]: plt.figure(figsize=(8, 8))

# Display 16 test images with predictions
for i in range(16):
    plt.subplot(4, 4, i+1)
    # Display each image
    plt.imshow(image_batch[i])
    # Set title color: green if prediction correct and red if prediction incorrect
    title_color = "green" if predicted_class_ids[i] == label_batch[i] else "red"
    plt.title(predicted_class_names[i].title(), color=title_color)
    plt.axis('off')

_ = plt.suptitle("Model Predictions")
```

## Model Predictions



[ ]: