

NAME-BACHU GANESH

ROLL NUMBER-CH.EN.U4AIE20003

---

### Question -1

Task - Handwritten Digit Classification

1. Build a CNN model to classify the digits.
2. Change the number of layers and check the performance of the model. Observe the results.
3. Try with different activation function and check the performance of the model. Observe the results.


```
In [1]: ▶ 1 import numpy as np
          2 import matplotlib.pyplot as plt
          3 import seaborn as sns
          4 import cv2
          5
          6 from PIL import Image
          7 import tensorflow as tf
          8 tf.random.set_seed(3)
          9 from tensorflow import keras
         10 from keras.datasets import mnist
         11 from tensorflow.math import confusion_matrix
```

```
In [2]: ▶ 1 (X_train, Y_train), (X_test, Y_test) = mnist.load_data()
```

```
In [3]: ▶ 1 # Shape of the numpy arrays
          2 print(X_train.shape, Y_train.shape, X_test.shape, Y_test.shape)

(60000, 28, 28) (60000,) (10000, 28, 28) (10000,)
```

```
In [4]: ▶ 1 # Displaying the image
          2 plt.imshow(X_train[25])
          3 plt.show()
```



```
In [5]: 1 # print the corresponding Label
        2 print(Y_train[25])
```

2

```
In [6]: 1 print(Y_train.shape, Y_test.shape)

        (60000,) (10000,)
```

```
In [7]: 1 # Unique values in Y_train
        2 print(np.unique(Y_train))
        3 # Unique values in Y_test
        4 print(np.unique(Y_test))
```

```
[0 1 2 3 4 5 6 7 8 9]
[0 1 2 3 4 5 6 7 8 9]
```

```
In [8]: 1 # Scaling the values
        2 #X_train = X_train/255
        3 #X_test = X_test/255
        4
        5 X_train = X_train.reshape(X_train.shape[0], 28, 28, 1)
        6 X_test = X_test.reshape(X_test.shape[0], 28, 28, 1)
        7 X_train, X_test = X_train / 255.0, X_test / 255.0
```

```
In [10]: 1 # Setting up the Layers of the Neural Network
        2 model = keras.Sequential([
        3     keras.layers.Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(28, 28,1)),
        4     keras.layers.MaxPooling2D(pool_size=(2, 2)),
        5     keras.layers.Flatten(),
        6     keras.layers.Dense(128, activation='relu'),
        7     keras.layers.Dense(64, activation='relu'),
        8     keras.layers.Dense(10, activation='softmax')])
```

```
In [11]: 1 # Compiling the Neural Network
        2 model.compile(optimizer='adam', loss = 'sparse_categorical_crossentropy', metrics=['accuracy'])
```

```
In [12]: 1 # Training the Neural Network
        2 model.fit(X_train, Y_train, epochs=2)
```


```
Epoch 1/2
1875/1875 [=====] - 53s 28ms/step - loss: 0.1569 - accuracy: 0.9531
Epoch 2/2
1875/1875 [=====] - 47s 25ms/step - loss: 0.0529 - accuracy: 0.9836
```

Out[12]: <keras.callbacks.History at 0x25c027195e0>

```
In [13]: 1 loss, accuracy = model.evaluate(X_test, Y_test)
        2 print(accuracy)
```

```
313/313 [=====] - 1s 4ms/step - loss: 0.0491 - accuracy: 0.9837
0.9836999773979187
```

```
In [14]: 1 # First data point in X_test
        2 plt.imshow(X_test[2])
        3 plt.show()
```



```
In [15]: 1 print(Y_test[2])
```

```
1
```

```
In [16]: 1 Y_pred = model.predict(X_test)
2 label_for_first_test_image = np.argmax(Y_pred[2])
3 print(label_for_first_test_image)

313/313 [=====] - 1s 4ms/step
1
```

```
In [17]: 1 Y_pred_label=np.argmax(Y_pred, axis=1)
2 print(Y_pred_label)

[7 2 1 ... 4 5 6]
```

```
In [18]: 1 #Confusion matrix
2 conf_mat = confusion_matrix(Y_test, Y_pred_label)
3 print(conf_mat)
```

```
tf.Tensor(
[[ 969  0  3  0  0  4  1  1  2  0]
 [  0 1123  5  1  1  1  2  0  2  0]
 [  0  1 1023  0  1  0  0  3  4  0]
 [  1  0  2  986  0 15  0  3  3  0]
 [  0  0  1  0  965  0  0  0  3 13]
 [  1  0  0  1  0  888  2  0  0  0]
 [  4  2  1  1  2  2  943  0  3  0]
 [  1  1 11  2  0  0  0 1010  2  1]
 [  5  0  4  0  1  2  1  4  954  3]
 [  4  2  1  1  4  7  0  5  9  976]], shape=(10, 10), dtype=int32)
```

After completing task 1 I came up with a conclusion that the model layers between the convolution layer (Input layer) and the output layer, the dense layers have increased the accuracy linear and the time taken to compile the model this not take much time with add extra dense layers but then add extra convolution layers the time taken to train the model has increased exponentially and there is no change in accuracy. The last layer must have sigmoid or SoftMax as activation function as we need the results between 0 to 1. There was no much different in accuracy or time taken when activation function as changed in last layer. In the other layers RELU has given the best accuracy when compared to tan-h and Leaky ReLU.

## Question-2

Task – Image Classification task

1. Build a CNN model to classify the different classes of images.
2. Change the training and testing split ratio. Observe the results.
3. Try with different cross validation approach, check the performance of the model. Observe the results.
4. Try removing max pooling and normalization layers, check the performance of the model. Observe the results.

I used a total of 3 different cross validations which are:-

1. Basic Cross Validation
2. K-Folds Cross Validation
3. Monte Carlo Cross Validation

### Basic Cross Validation

```
In [13]: 1 import tensorflow as tf
2 from tensorflow.keras import datasets, layers, models
3 import matplotlib.pyplot as plt
4 import numpy as np
5 from sklearn.model_selection import cross_val_score
6 from sklearn.model_selection import train_test_split
7
8 from keras.datasets import cifar10
9 from sklearn.cross_validation import train_test_split
10 from sklearn.metrics import classification_report, confusion_matrix
```

```
In [14]: 1 (X_train, Y_train), (X_test, Y_test) = cifar10.load_data()
2
3 # View first image
4 plt.imshow(X_train[0])
5 plt.show()
6
7 import numpy as np
8 from sklearn.model_selection import train_test_split
9
10 # Concatenate train and test images
11 X = np.concatenate((X_train,X_test))
12 y = np.concatenate((Y_train,Y_test))
13
14 # Check shape
15 print(X.shape) # (60000, 32, 32, 3)
16
17 # Split data
18 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=10000, random_state=1234)
19
20 # Check shape
21 print(X_train.shape) # (50000, 32, 32, 3)
22
23 # View first image
24 plt.imshow(X_train[0])
25 plt.show()
```

```
In [15]: 1 X_test.shape
```

```
Out[15]: (10000, 32, 32, 3)
```

```
In [16]: 1 y_train.shape
```

```
Out[16]: (50000, 1)
```

```
In [17]: 1 y_train[:5]
```

```
Out[17]: array([[4],
                [0],
                [1],
                [3],
                [8]], dtype=uint8)
```

```
In [18]: 1 y_train = y_train.reshape(-1,)
          2 y_train[:5]
```


```
Out[18]: array([4, 0, 1, 3, 8], dtype=uint8)
```

```
In [19]: 1 y_test = y_test.reshape(-1,)
```


```
In [20]: 1 classes = ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog", "horse", "ship", "truck"]
```

```
In [21]: 1 def plot_sample(X, y, index):
          2     plt.figure(figsize = (15,2))
          3     plt.imshow(X[index])
          4     plt.xlabel(classes[y[index]])
```

```
In [22]: 1 plot_sample(X_train, y_train, 0)
```



```
In [23]: 1 plot_sample(X_train, y_train, 1)
```



```
In [24]: 1 X_train = X_train / 255.0
          2 X_test = X_test / 255.0
```

```
In [25]: 1 cnn = models.Sequential([
2         layers.Conv2D(filters=32, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)),
3         layers.MaxPooling2D((2, 2)),
4         layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu'),
5         layers.MaxPooling2D((2, 2)),
6         layers.Flatten(),
7         layers.Dense(64, activation='relu'),
8         layers.Dense(10, activation='softmax')
9     ])
```

```
In [26]: 1 cnn.compile(optimizer='adam',
2         loss='sparse_categorical_crossentropy',
3         metrics=['accuracy'])
```

```
In [30]: 1 #cnn.fit(X_train, y_train, epochs=2)
2
3 n_folds=3
4 epochs=20
5 batch_size=128
6
7 #save the model history in a list after fitting so that we can plot later
8 model_history = []
9
10 for i in range(n_folds):
11     print("Training on Fold: ", i+1)
12     cnn.fit(X_train, y_train, epochs=2)
13     print("====="*12, end="\n\n\n")
```

```
Training on Fold:  1
Epoch 1/2
1563/1563 [=====] - 59s 38ms/step - loss: 1.0039 - accuracy: 0.6495
Epoch 2/2
1563/1563 [=====] - 65s 41ms/step - loss: 0.9305 - accuracy: 0.6778
=====
```

```
Training on Fold:  2
Epoch 1/2
1563/1563 [=====] - 64s 41ms/step - loss: 0.8639 - accuracy: 0.6999
Epoch 2/2
1563/1563 [=====] - 57s 37ms/step - loss: 0.8124 - accuracy: 0.7165
=====
```

```
Training on Fold:  3
Epoch 1/2
1563/1563 [=====] - 57s 36ms/step - loss: 0.7666 - accuracy: 0.7350
Epoch 2/2
1563/1563 [=====] - 68s 44ms/step - loss: 0.7168 - accuracy: 0.7508
=====
```

```
In [31]: 1 cnn.evaluate(X_test, y_test)

313/313 [=====] - 5s 13ms/step - loss: 0.9167 - accuracy: 0.6939
```

```
Out[31]: [0.9167031049728394, 0.6938999891281128]
```

```
In [32]: 1 y_pred = cnn.predict(X_test)
2         y_pred[:5]
3
4         clas=np.argmax(y_pred,axis=1)
5         clas

313/313 [=====] - 4s 13ms/step
```

```
Out[32]: array([8, 5, 2, ..., 2, 6, 3], dtype=int64)
```


```
In [33]: 1 y_classes = [np.argmax(element) for element in y_pred]
        2 y_classes[:5]
```

```
Out[33]: [8, 5, 2, 7, 8]
```

```
In [34]: 1 y_test[:5]
```

```
Out[34]: array([8, 3, 3, 7, 8], dtype=uint8)
```

```
In [35]: 1 plot_sample(X_test, y_test, 3)
```



```
In [36]: 1 classes[y_classes[3]]
```

```
Out[36]: 'horse'
```

## K-Folds Cross Validation

```
In [1]: 1 import numpy as np
        2 import tensorflow as tf
        3 from tensorflow import keras
        4 from sklearn.model_selection import StratifiedKFold
```

```
In [2]: 1 # Load the CIFAR10 dataset
        2 (x_train, y_train), (x_test, y_test) = keras.datasets.cifar10.load_data()
        3
        4 # Normalize the data
        5 x_train = x_train.astype('float32') / 255
        6 x_test = x_test.astype('float32') / 255
        7
```

```
In [3]: 1 # Define the model architecture
        2 model = keras.Sequential([
        3     keras.layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
        4     keras.layers.MaxPooling2D((2, 2)),
        5     keras.layers.Conv2D(64, (3, 3), activation='relu'),
        6     keras.layers.MaxPooling2D((2, 2)),
        7     keras.layers.Flatten(),
        8     keras.layers.Dense(128, activation='relu'),
        9     keras.layers.Dense(10)
        10 ])
```

```
In [4]: 1 # Compile the model
        2 model.compile(optimizer='adam', loss=tf.losses.SparseCategoricalCrossentropy(from_logits=True), metrics=['accuracy'])
```

```
In [5]: 1 # Define the number of folds for k-fold cross validation
        2 k = 5
        3
        4 # Define the cross validation iterator
        5 skf = StratifiedKFold(n_splits=k, shuffle=True, random_state=1)
        6
        7 # Initialize an array to store the accuracy scores for each fold
        8 scores = []
```





```

In [10]: # Define the number of iterations for Monte Carlo cross validation
n_iter = 5

# Define the cross validation iterator
shuffle_split = ShuffleSplit(n_splits=n_iter, test_size=0.3, random_state=1)

# Initialize an array to store the accuracy scores for each iteration
scores = []

# Loop over the iterations
for train_index, val_index in shuffle_split.split(x_train, y_train):
    # Split the data into training and validation sets for the current iteration
    x_train_mc, x_val_mc = x_train[train_index], x_train[val_index]
    y_train_mc, y_val_mc = y_train[train_index], y_train[val_index]

    # Train the model on the current iteration
    model.fit(x_train_mc, y_train_mc, batch_size=32, epochs=10, validation_data=(x_val_mc, y_val_mc))

    # Evaluate the model on the validation set for the current iteration
    val_loss, val_accuracy = model.evaluate(x_val_mc, y_val_mc)

    # Append the accuracy score for the current iteration to the scores array
    scores.append(val_accuracy)

1094/1094 [=====] - 21s 19ms/step - loss: 0.3137 - accuracy: 0.9116 - val_loss: 0.2893 - val_accuac
y: 0.9179

In [11]: # Print the mean accuracy and standard deviation of the accuracy scores
print("Accuracy: %.2f%% (+/- %.2f%%)" % (np.mean(scores) * 100, np.std(scores) * 100))

# Evaluate the model on the test set
test_loss, test_accuracy = model.evaluate(x_test, y_test)
print("Test Accuracy: %.2f%%" % (test_accuracy * 100))

Accuracy: 78.11% (+/- 7.37%)
313/313 [=====] - 3s 9ms/step - loss: 3.0424 - accuracy: 0.6675
Test Accuracy: 66.75%

```

On removing the Max pooling layer, the time required to compile the models has reduced and the accuracy also dropped this is due to over fitting in the model. On removing the normalization layer, the training has increased and the accuracy has dropped this is due to the large number the model has to compute. On increasing the training dataset size the accuracy has increased but this had lead to over fitting in the model and the training time has increase linearly.